

The Influence of Fatigue on Usage of Model-Based vs Model-Free Reinforcement Learning Strategies

Colton Loftus, Maya Rozenshteyn

2022-12-15

Contents

Tiredness Plotting Method 1: Tiredness Scores 13

Tiredness Plotting Method 2: Relative Tiredness between Trials 15

```
setwd("./")
```

```
library(purrr)
library(readr)
library(jsonlite)
```

```
Attaching package: 'jsonlite'
The following object is masked from 'package:purrr':
```

```
    flatten
library("data.table")
```

```
Attaching package: 'data.table'
The following object is masked from 'package:purrr':
```

```
    transpose
The following objects are masked from 'package:dplyr':
```

```
    between, first, last
library(dplyr)
library(rlist)
library(magrittr)
```

```
Attaching package: 'magrittr'
The following object is masked from 'package:purrr':
```

```
    set_names
library(lme4)
Loading required package: Matrix
library(plotrix)
library(ggplot2)
```

```
# Preprocess experiment data to extract relevant information: reward,
# transition, first state choice, tiredness, experimental condition (morning
# vs evening), and task run (first or second time conducting the experiment).
```

```

# Note 1: reward is coded as -1 (no reward) and 1 (reward), transition is coded as
# --0.7 (rare transition) and ~0.3 (common transition), tiredness is code on a
# scale from -3 (least tired) to 3 (most tired), experimental condition is coded
# as 1 (morning) and -1 (evening) and task run as -1 (first run) and 1
# (second run).
# Note 2: "stay" is true when the participant selected the same first-stage
# choice in successive trials.
# Note 3: all predictors are zero-centered.
survey_response_to_val <- jsonlite::fromJSON("data/sleepiness.json")$meta
preprocess <- function(csv, survey_string, is_morning, morning_first, is_more_tired_condition, person_id) {
  relevant_data <- csv %>%
    dplyr::filter(trial_type == "two-step-trial") %>%
    select(state_1_choice, transition, outcome) %>%
    dplyr::filter(!grepl("null", state_1_choice)) %>%
    dplyr::filter(!grepl("null", transition)) %>%
    dplyr::filter(!grepl("null", outcome)) %>%
    mutate(state_1_choice = as.numeric(state_1_choice)) %>%
    mutate(outcome = as.numeric(outcome)) %>%
    mutate(next_state1_choice = lead(state_1_choice, 1)) %>%
    mutate(stay = state_1_choice == next_state1_choice) %>%
    mutate(tiredness = survey_response_to_val[survey_string][[1]]) %>%
    mutate(is_morning = is_morning) %>%
    mutate(task_run = ((is_morning && morning_first) || (is_morning && morning_first))) %>%
    mutate(is_morning = ifelse(is_morning == TRUE, 1, -1)) %>%
    mutate(task_run = ifelse(task_run == TRUE, -1, 1)) %>%
    mutate(subject = person_id) %>%
    dplyr::filter(!is.na(next_state1_choice)) %>%
    mutate(outcome = ifelse(outcome == 0, -1, outcome)) %>%
    as.data.frame()

  attr(relevant_data, "is_morning") <- is_morning
  attr(relevant_data, "task_run") <- (is_morning && morning_first) || (is_morning && morning_first)
  attr(relevant_data, "tiredness") <- survey_response_to_val[survey_string][[1]]
  attr(relevant_data, "subject") <- person_id
  attr(relevant_data, "is_more_tired_condition") <- is_more_tired_condition

  return(relevant_data)
}

```

```

all_data <- list()
# Read in and preprocess experimental data.
for (person in list.files("data/")) {
  if (!grepl("json", person)) {
    survey_response_map <- jsonlite::fromJSON(paste0("data/", person, "/meta.json"))$meta
    morning_first <- survey_response_map$first == "morning"
    for (csv in list.files(paste0("data/", person, "/"))) {
      if (!grepl("json", csv)) {
        csv_is_morning <- grepl("day", csv)
        if (csv_is_morning && morning_first | !csv_is_morning && !morning_first) {
          survey_response <- survey_response_map["form1"]
          is_more_tired_condition <- (survey_response_map["form1"][[1]] >= survey_response_map["form1"][[2]])
        } else {
          survey_response <- survey_response_map["form2"]
          is_more_tired_condition <- (survey_response_map["form2"][[1]] > survey_response_map["form2"][[2]])
        }
      }
    }
  }
}

```

```

    }
    read_in_csv <- read_csv(paste0("data/", person, "/", csv))
    processed_csv <- preprocess(read_in_csv, survey_response[[1]], csv_is_morning, morning_
    all_data[[length(all_data) + 1]] <- processed_csv
  }
}
}

all_data_concat <- bind_rows(all_data) %>% mutate(transition = transition - mean(transition))

# Mixed effect modeling illustrating random and fixed effects of outcome and
# transition on stay.
base_reg <- glmer(stay ~ outcome*transition + (1 | subject), data = all_data_concat, family = binomial)
summary(base_reg)
Generalized linear mixed model fit by maximum likelihood
(Laplace Approximation) [glmerMod]
Family: binomial (logit)
Formula: stay ~ outcome * transition + (1 | subject)
Data: all_data_concat

      AIC      BIC   logLik deviance df.resid
4236.1  4268.6  -2113.0  4226.1     4973

Scaled residuals:
    Min       1Q   Median       3Q      Max
-7.8028  0.1570  0.3464  0.4352  1.2803

Random effects:
Groups Name      Variance Std.Dev.
subject (Intercept) 1.052    1.026
Number of obs: 4978, groups: subject, 16

Fixed effects:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)    1.73443    0.26089   6.648 2.97e-11 ***
outcome         0.12957    0.03979   3.257  0.00113 **
transition     -0.11357    0.08795  -1.291  0.19663
outcome:transition 0.24767    0.08805   2.813  0.00491 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:
      (Intr) outcom trnstn
outcome    0.028
transition -0.005  0.009
otcm:trnstn 0.005 -0.021  0.153

# Mixed effect modeling illustrating random and fixed effects of outcome,
# transition, task run, and time of day on stay.
tod_reg <- glmer(stay ~ outcome*transition*(task_run + is_morning) + (1 | subject), data = all_data_concat)
summary(tod_reg)
Generalized linear mixed model fit by maximum likelihood
(Laplace Approximation) [glmerMod]

```

```

Family: binomial ( logit )
Formula:
stay ~ outcome * transition * (task_run + is_morning) + (1 |
  subject)
Data: all_data_concat

      AIC      BIC  logLik deviance df.resid
4234.7  4319.3 -2104.3  4208.7    4965

Scaled residuals:
      Min       1Q   Median       3Q      Max
-7.9662  0.1630  0.3540  0.4498  1.3296

Random effects:
Groups Name      Variance Std.Dev.
subject (Intercept) 1.069    1.034
Number of obs: 4978, groups: subject, 16

Fixed effects:

```

	Estimate	Std. Error	z value
(Intercept)	1.65746	0.26553	6.242
outcome	0.19214	0.04840	3.970
transition	-0.06438	0.10644	-0.605
task_run	0.17380	0.07858	2.212
is_morning	0.13738	0.05631	2.440
outcome:transition	0.20535	0.10651	1.928
outcome:task_run	-0.14010	0.05767	-2.429
outcome:is_morning	-0.04169	0.05028	-0.829
transition:task_run	-0.12928	0.12809	-1.009
transition:is_morning	-0.03410	0.11246	-0.303
outcome:transition:task_run	0.10996	0.12812	0.858
outcome:transition:is_morning	-0.02721	0.11248	-0.242

```

Pr(>|z|)
(Intercept) 4.32e-10 ***
outcome      7.19e-05 ***
transition   0.5453
task_run     0.0270 *
is_morning   0.0147 *
outcome:transition 0.0538 .
outcome:task_run 0.0151 *
outcome:is_morning 0.4070
transition:task_run 0.3128
transition:is_morning 0.7617
outcome:transition:task_run 0.3908
outcome:transition:is_morning 0.8089
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:
      (Intr) outcom trnstn tsk_rn is_mrn otcm:t
outcome      0.045
transition   -0.003 -0.007
task_run     -0.139 -0.114 -0.006

```

```

is_morning      -0.094 -0.048  0.006  0.706
otcm:trnstn      0.002 -0.012  0.217  0.007 -0.006
otcm:tsk_rn     -0.029 -0.562  0.010  0.123  0.049  0.003
otcm:s_mrnn     -0.012 -0.317 -0.003  0.049  0.103  0.009
trnstn:tsk_      0.001  0.010 -0.554 -0.027 -0.035 -0.147
trnstn:s_mr      0.001 -0.005 -0.316 -0.021 -0.045 -0.039
otcm:trnstn:t_   0.001  0.001 -0.147  0.019  0.035 -0.553
otcm:trnstn:s_  -0.003  0.010 -0.038  0.028  0.044 -0.316
               otc_m:t_ otc_m:s_ trnstn:t_ trnstn:s_

```

```

outcome
transition
task_run
is_morning
otcm:trnstn
otcm:tsk_rn
otcm:s_mrnn      0.607
trnstn:tsk_      0.025  0.036
trnstn:s_mr      0.036  0.047  0.616
otcm:trnstn:t_  -0.035 -0.036  0.138  0.017
otcm:trnstn:s_  -0.037 -0.054  0.018  0.057
               otc_m:trnstn:t_

```

```

outcome
transition
task_run
is_morning
otcm:trnstn
otcm:tsk_rn
otcm:s_mrnn
trnstn:tsk_
trnstn:s_mr
otcm:trnstn:t_
otcm:trnstn:s_  0.616

```

```

# Mixed effect modeling illustrating random and fixed effects of outcome,
# transition, task run, and tiredness on stay.

```

```

tiredness_reg <- glmer(stay ~ outcome*transition*(task_run + tiredness) + (1 | subject), data = all_data)
summary(tiredness_reg)

```

```

Generalized linear mixed model fit by maximum likelihood

```

```

(Laplace Approximation) [glmerMod]

```

```

Family: binomial (logit)

```

```

Formula:

```

```

stay ~ outcome * transition * (task_run + tiredness) + (1 | subject)

```

```

Data: all_data_concat

```

```

      AIC      BIC   logLik deviance df.resid
4207.0   4291.6  -2090.5   4181.0     4965

```

```

Scaled residuals:

```

```

      Min       1Q   Median       3Q      Max
-8.6816  0.1629  0.3387  0.4576  1.2465

```

```

Random effects:

```

```

Groups  Name      Variance Std.Dev.
subject (Intercept) 1.021    1.01

```

Number of obs: 4978, groups: subject, 16

Fixed effects:

	Estimate	Std. Error	z	value
(Intercept)	1.90075	0.27577	6.893	
outcome	0.44841	0.07400	6.060	
transition	0.23751	0.15229	1.560	
task_run	0.02491	0.05631	0.442	
tiredness	0.09960	0.06194	1.608	
outcome:transition	0.57259	0.15258	3.753	
outcome:task_run	-0.16694	0.04775	-3.496	
outcome:tiredness	0.14152	0.02996	4.723	
transition:task_run	-0.16360	0.10405	-1.572	
transition:tiredness	0.15217	0.06218	2.447	
outcome:transition:task_run	0.04690	0.10406	0.451	
outcome:transition:tiredness	0.19969	0.06244	3.198	

Pr(>|z|)

(Intercept)	5.48e-12	***
outcome	1.36e-09	***
transition	0.118862	
task_run	0.658304	
tiredness	0.107838	
outcome:transition	0.000175	***
outcome:task_run	0.000472	***
outcome:tiredness	2.32e-06	***
transition:task_run	0.115871	
transition:tiredness	0.014391	*
outcome:transition:task_run	0.652217	
outcome:transition:tiredness	0.001384	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	outcom	trnstn	tsk_rn	trdnss	otcm:trn	otcm:_
outcome	0.107						
transition	0.037	0.208					
task_run	-0.136	-0.113	-0.026				
tiredness	0.345	0.162	0.063	-0.110			
otcm:trnstn	0.057	0.141	0.257	-0.015	0.107		
otcm:tsk_rn	-0.044	-0.491	-0.036	0.191	-0.027	-0.031	
otcm:trdnss	0.090	0.777	0.205	-0.024	0.180	0.135	-0.256
trnstn:tsk_	-0.012	-0.035	-0.494	0.002	-0.021	-0.144	0.024
trnstn:trdn	0.032	0.201	0.747	-0.012	0.044	0.169	-0.034
otcm:trns:_	-0.010	-0.031	-0.143	0.011	-0.026	-0.490	0.006
otcm:trnst:	0.055	0.131	0.169	-0.017	0.098	0.748	-0.014
otcm:trd trns:_							
otcm:trnst:							
otcm:trdnss							

```

trnstn:tsk_ -0.035
trnstn:trdn  0.188  -0.255
otcm:trns:_ -0.015   0.225 -0.029
otcm:trnst:  0.119  -0.029  0.179 -0.251

sum_decision_type <- function(df, stay, outcome, transition) {
  return(sum(df$stay == stay & df$outcome == outcome & df$transition == transition))
}

stay_probability_computer <- function(df) {
  # Common Transition, Rewarded, Stay
  ct_r_s <- sum_decision_type(df, stay = TRUE, outcome = 1, transition = 1)

  # Common Transition, Rewarded, Leave
  ct_r_l <- sum_decision_type(df, stay = FALSE, outcome = 1, transition = 1)

  # Rare Transition, Rewarded, Stay
  rt_r_s <- sum_decision_type(df, stay = TRUE, outcome = 1, transition = 0)

  # Rare Transition, Rewarded, Leave
  rt_r_l <- sum_decision_type(df, stay = FALSE, outcome = 1, transition = 0)

  # Common Transition, Unrewarded, Stay
  ct_ur_s <- sum_decision_type(df, stay = TRUE, outcome = -1, transition = 1)

  # Common Transition, Unrewarded, Leave
  ct_ur_l <- sum_decision_type(df, stay = FALSE, outcome = -1, transition = 1)

  # Rare Transition, Unrewarded, Stay
  rt_ur_s <- sum_decision_type(df, stay = TRUE, outcome = -1, transition = 0)

  # Rare Transition, Unrewarded, Leave
  rt_ur_l <- sum_decision_type(df, stay = FALSE, outcome = -1, transition = 0)

  # Common Transition, Rewarded
  b1 <- ct_r_s / (ct_r_s + ct_r_l)
  # Rare Transition, Rewarded
  b2 <- rt_r_s / (rt_r_s + rt_r_l)
  # Common Transition, Unrewarded
  b3 <- ct_ur_s / (ct_ur_s + ct_ur_l)
  # Rare Transition, Unrewarded
  b4 <- rt_ur_s / (rt_ur_s + rt_ur_l)

  res = c(b1, b2, b3, b4)
  names(res) <- c("b1", "b2", "b3", "b4")
  return(res)
}

bp_val_averger <- function(bp_vals) {
  b1s <- c()
  b2s <- c()
  b3s <- c()
  b4s <- c()
  for (i in 1:length(bp_vals)) {

```

```

    b1s <- c(b1s, bp_vals[[i]][["b1"]])
    b2s <- c(b2s, bp_vals[[i]][["b2"]])
    b3s <- c(b3s, bp_vals[[i]][["b3"]])
    b4s <- c(b4s, bp_vals[[i]][["b4"]])
  }

  b1_av <- mean(b1s)
  b2_av <- mean(b2s)
  b3_av <- mean(b3s)
  b4_av <- mean(b4s)

  b1_stderr <- std.error(b1s)
  b2_stderr <- std.error(b2s)
  b3_stderr <- std.error(b3s)
  b4_stderr <- std.error(b4s)

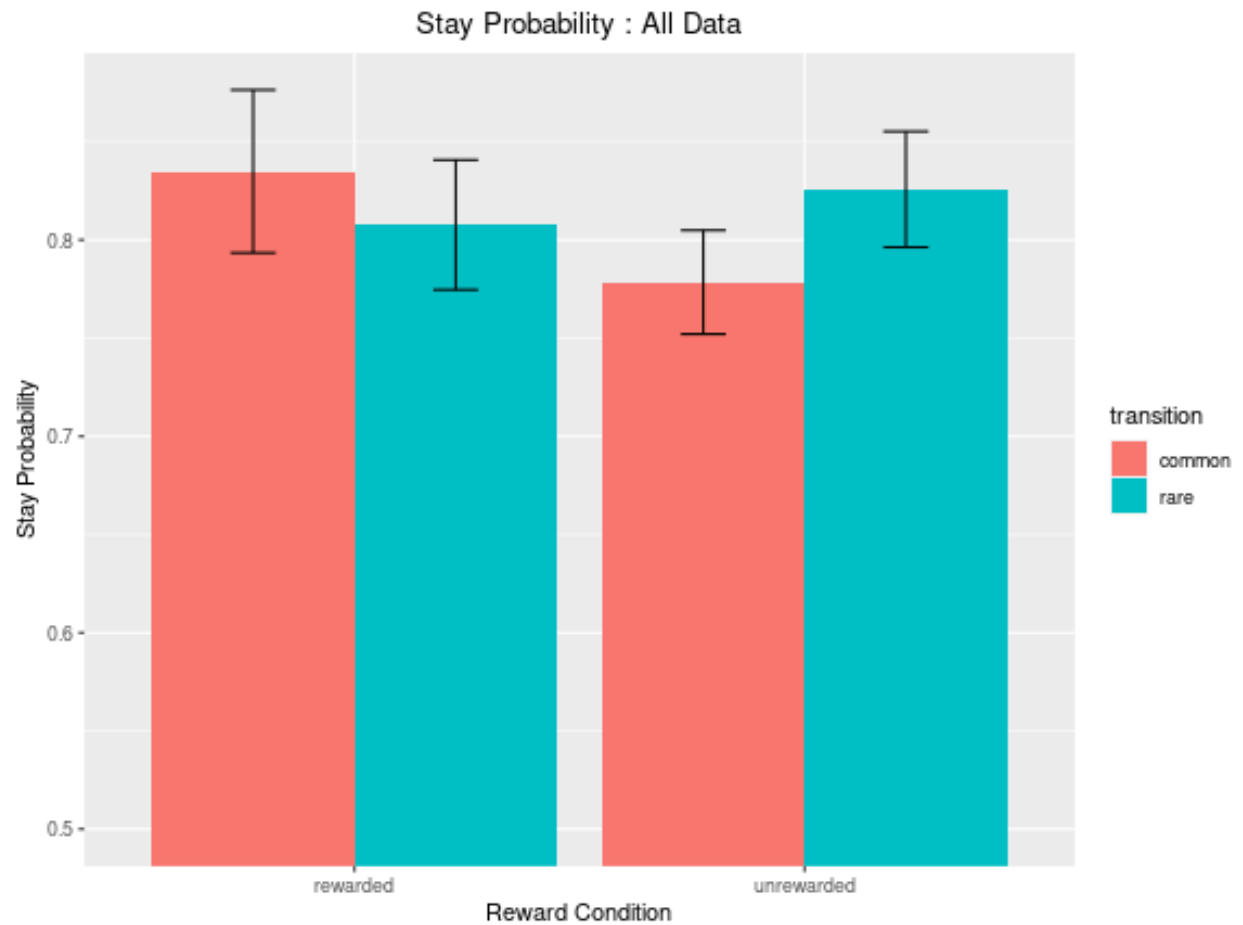
  return(c(b1_av, b2_av, b3_av, b4_av, b1_stderr, b2_stderr, b3_stderr, b4_stderr))
}

stay_bar_plot_creator <- function(bars_to_plot, title) {
  reward_type <- c(rep("rewarded", 2), rep("unrewarded", 2))
  transition <- rep(c("common", "rare"), 2)
  value <- bars_to_plot[1:4]
  se <- bars_to_plot[5:8]
  data <- data.frame(reward_type, transition, value)
  ggplot(data, aes(fill = transition, y = value, x = reward_type)) +
    geom_bar(position = "dodge", stat = "identity") +
    labs(y= "Stay Probability", x = "Reward Condition", title = title) +
    coord_cartesian(ylim = c(.5, max(value) + max(se))) +
    theme(plot.title = element_text(hjust = 0.5)) +
    geom_errorbar(aes(ymin=value-se, ymax=value+se),
                  width=.2,
                  position=position_dodge(.9))
}


all_data_processed <- all_data %>%
  map(stay_probability_computer)

stay_bar_plot_creator(bp_val_averger(all_data_processed), "Stay Probability : All Data")


```

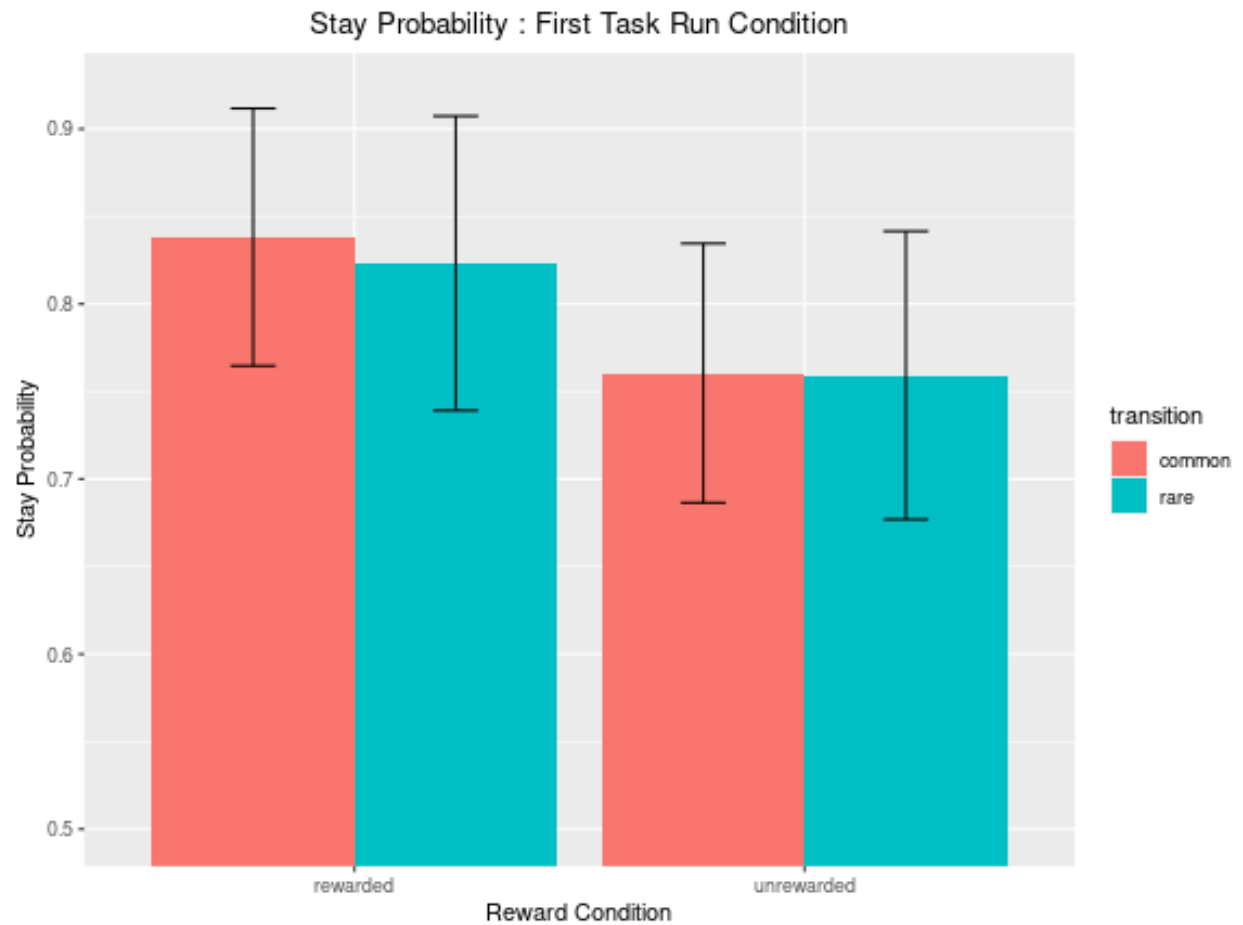
```
morning_data <- all_data %>%  
  list.filter(attr(., "is_morning") == TRUE) %>%  
  map(stay_probability_computer)  
  
stay_bar_plot_creator(bp_val_averger(morning_data), "Stay Probability : Morning Condition")
```



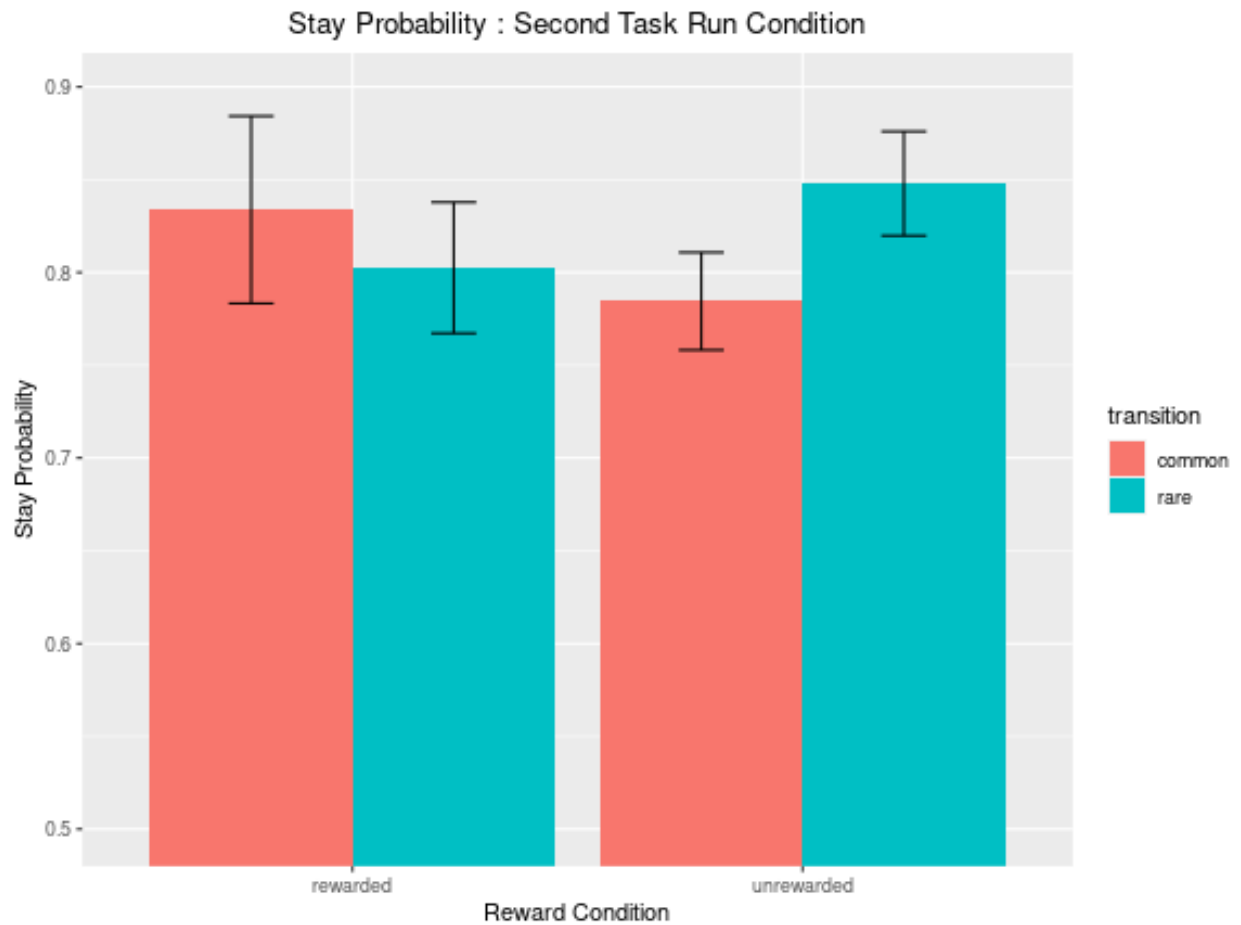
```
evening_data <- all_data %>%  
  list.filter(attr(., "is_morning") == FALSE) %>%  
  map(stay_probability_computer)  
  
stay_bar_plot_creator(bp_val_averger(evening_data), "Stay Probability : Evening Condition")
```



```
first_run_data <- all_data %>%  
  list.filter(attr(., "task_run") == TRUE) %>%  
  map(stay_probability_computer)  
  
stay_bar_plot_creator(bp_val_averger(first_run_data), "Stay Probability : First Task Run Condition")
```

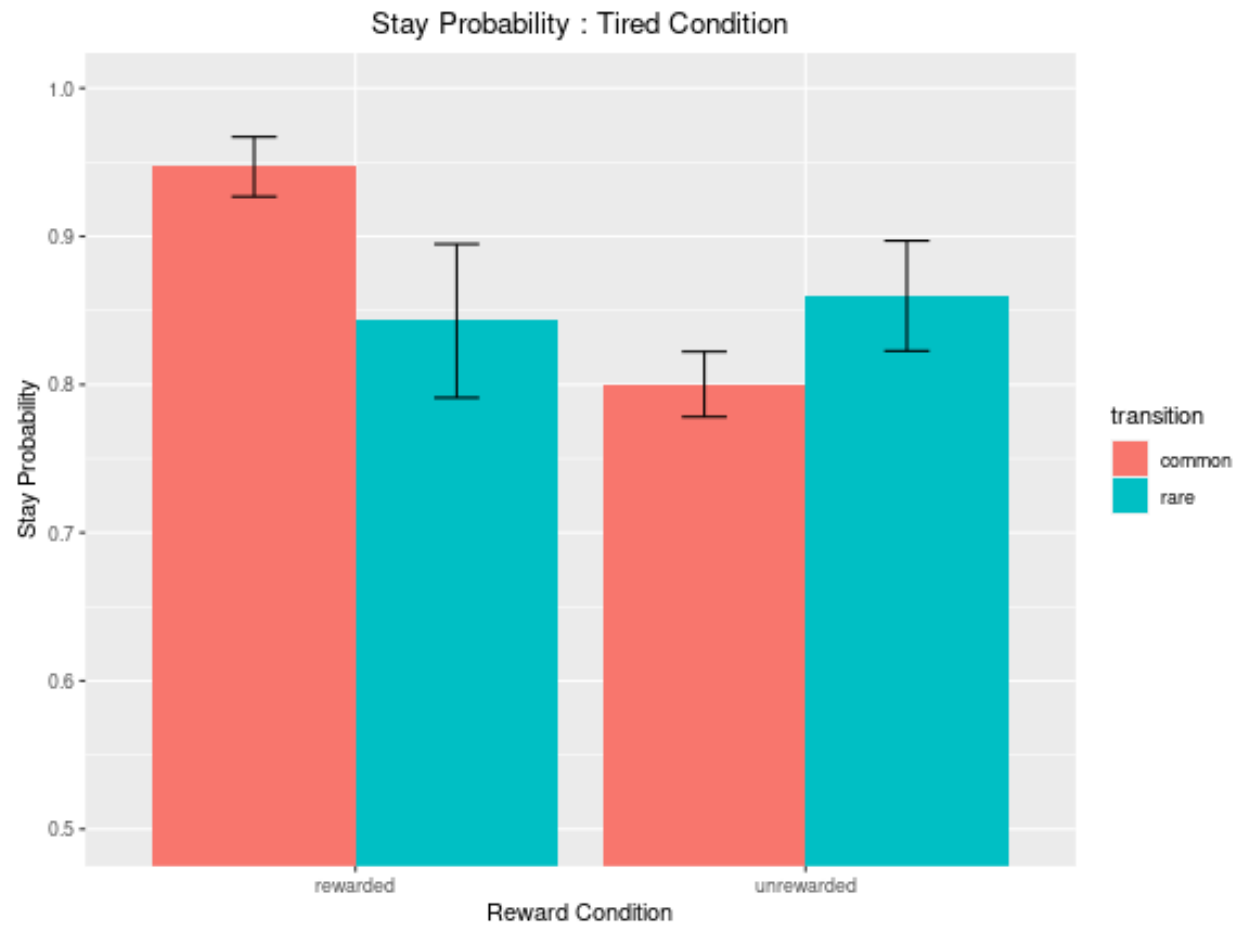


```
second_run_data <- all_data %>%  
  list.filter(attr(., "task_run") == FALSE) %>%  
  map(stay_probability_computer)  
  
stay_bar_plot_creator(bp_val_averger(second_run_data), "Stay Probability : Second Task Run Condition")
```

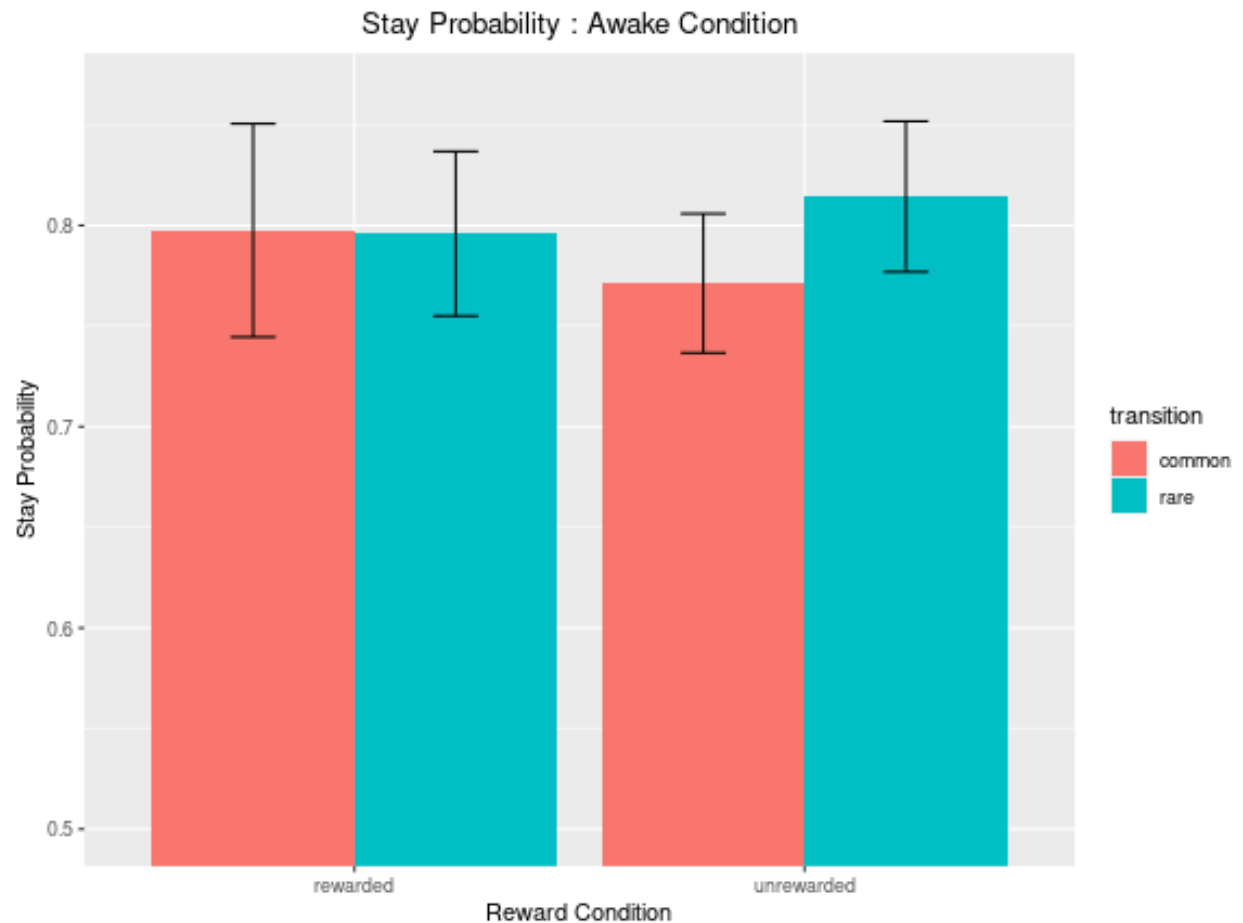


Tiredness Plotting Method 1: Tiredness Scores

```
tired_data <- all_data %>%  
  list.filter(attr(., "tiredness") >= 0) %>%  
  map(stay_probability_computer)  
  
stay_bar_plot_creator(bp_val_averger(tired_data), "Stay Probability : Tired Condition")
```




```
awake_data <- all_data %>%  
  list.filter(attr(., "tiredness") < 0) %>%  
  map(stay_probability_computer)  
  
stay_bar_plot_creator(bp_val_averger(awake_data), "Stay Probability : Awake Condition")
```



Tiredness Plotting Method 2: Relative Tiredness between Trials

```
relatively_tired_data <- all_data %>%
  list.filter(attr(., "is_more_tired_condition") == TRUE) %>%
  map(stay_probability_computer)

stay_bar_plot_creator(bp_val_averger(relatively_tired_data), "Stay Probability : Tired Condition")
```



```
relatively_awake_data <- all_data %>%  
  list.filter(attr(., "is_more_tired_condition") == FALSE) %>%  
  map(stay_probability_computer)  
  
stay_bar_plot_creator(bp_val_averger(relatively_awake_data), "Stay Probability : Awake Condition")
```